START HERE

One giant prompt will fail your workflow

We tried using a single prompt to research, draft, edit, and design content. It failed consistently. Here is what we learned from rebuilding it.

READ THIS

Split the work into specialized agents

We replaced our single prompt with seven task-specific agents. One drafts, one checks citations, one guards tone, and one formats text. Accuracy improved the moment each agent received one job.

READ THIS

Match models to operational needs

Heavy reasoning tasks run on Gemini 3.1 Pro. Fast, repetitive tasks run on Gemini 3.8 Flash. Assigning models based on role keeps costs down without dropping quality.

READ THIS

Build automatic retries and failovers

Endpoints flake during long generations. If our social copy agent fails, it retries twice before failing over to a larger model. The screen never displays half-written data.

READ THIS

Let code handle the visual design

Brand consistency drifts when models manage layout. Our agents return raw text in structured JSON. A deterministic code renderer maps that text directly into fixed brand templates.

READ THIS

Shift the operator to reviewer

Automation handles drafting and layout. Every final asset drops into an approval queue. The human operator never writes from scratch, but always gives the final sign-off.

THE CLOSE

Keep your system modular

Save this post for the next time you design an automated content pipeline.